2CS702 Big Data Analytics

Lab-3 Task

Submitted by: Labdhi Sheth 18BCE101

Aim: **Setup single-node Hadoop cluster and apply HDFS commands on single-node Hadoop Cluster.**

Methodology followed:
The installation process as shared in the drive was followed.

Output:
CMD run as administrator

1. Checking the java and Hadoop version

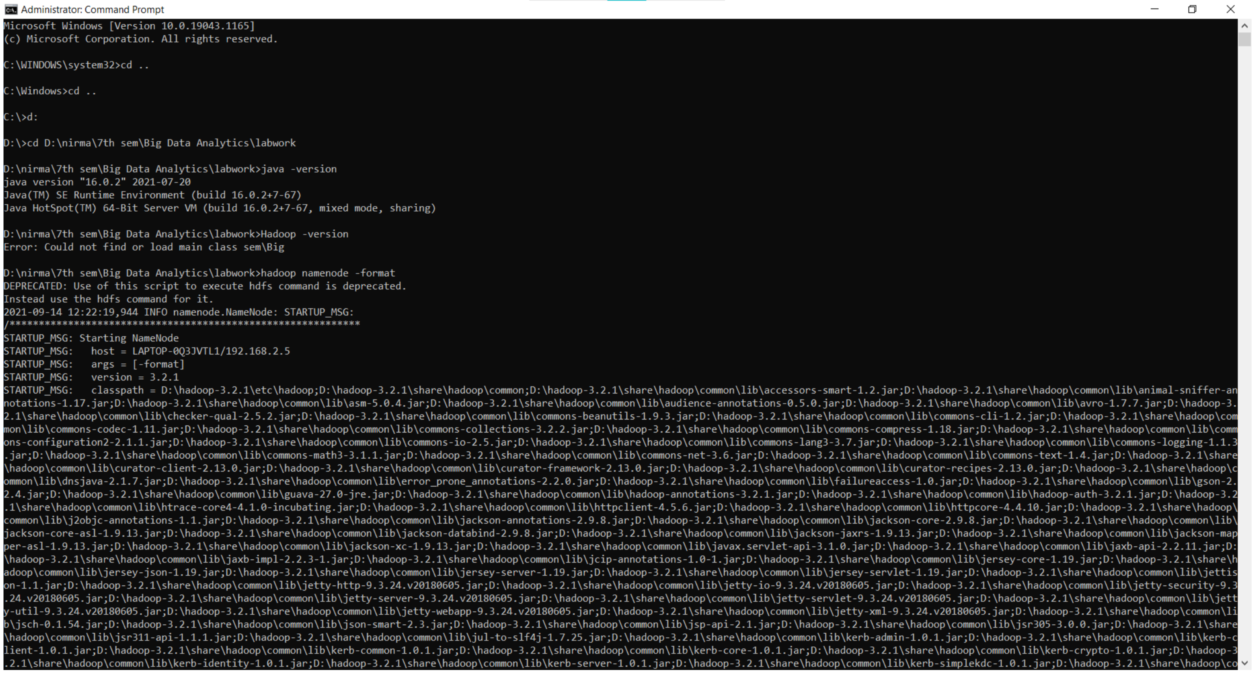

Administrator: Command Prompt

```
Microsoft Windows [Version 10.0.19043.1165]
(c) Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>cd ..

C:\Windows>cd ..

C:\>d:

D:\>cd D:\nirma\7th sem\Big Data Analytics\labwork

D:\nirma\7th sem\Big Data Analytics\labwork>java -version
java version "16.0.2" 2021-07-20
Java(TM) SE Runtime Environment (build 16.0.2+7-67)
Java HotSpot(TM) 64-Bit Server VM (build 16.0.2+7-67, mixed mode, sharing)

D:\nirma\7th sem\Big Data Analytics\labwork>Hadoop -version
Error: Could not find or load main class sem\Big

D:\nirma\7th sem\Big Data Analytics\labwork>hadoop namenode -format
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
2021-09-14 12:22:19,944 INFO namenode.NameNode: STARTUP_MSG:
/************************************************************
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = LAPTOP-0Q3JVTL1/192.168.2.5
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 3.2.1
STARTUP_MSG:   classpath = D:\hadoop-3.2.1\etc\hadoop;D:\hadoop-3.2.1\share\hadoop\common;D:\hadoop-3.2.1\share\hadoop\common\lib\accessors-smart-1.2.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\animal-sniffer-an
notations-1.17.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\asm-5.0.4.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\audience-annotations-0.5.0.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\avro-1.7.7.jar;D:\hadoop-3.
2.1\share\hadoop\common\lib\checker-qual-2.5.2.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-beanutils-1.9.3.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-cli-1.2.jar;D:\hadoop-3.2.1\share\hadoop\com
mon\lib\commons-codec-1.11.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-collections-3.2.2.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-compress-1.18.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\comm
ons-configuration2-2.1.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-io-2.5.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-lang3-3.7.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-logging-1.1.3
.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-math3-3.1.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-net-3.6.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\commons-text-1.4.jar;D:\hadoop-3.2.1\share
\hadoop\common\lib\curator-client-2.13.0.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\curator-framework-2.13.0.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\curator-recipes-2.13.0.jar;D:\hadoop-3.2.1\share\hadoop\c
ommon\lib\dnsjava-2.1.7.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\error_prone_annotations-2.2.0.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\failureaccess-1.0.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\gson-2.
2.4.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\guava-27.0-jre.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\hadoop-annotations-3.2.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\hadoop-auth-3.2.1.jar;D:\hadoop-3.2
.1\share\hadoop\common\lib\htrace-core4-4.1.0-incubating.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\httpclient-4.5.6.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\httpcore-4.4.10.jar;D:\hadoop-3.2.1\share\hadoop\
common\lib\j2objc-annotations-1.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jackson-annotations-2.9.8.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jackson-core-2.9.8.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\
jackson-core-asl-1.9.13.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jackson-databind-2.9.8.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jackson-jaxrs-1.9.13.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jackson-map
per-asl-1.9.13.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jackson-xc-1.9.13.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\javax.servlet-api-3.1.0.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jaxb-api-2.2.11.jar;D:
\hadoop-3.2.1\share\hadoop\common\lib\jaxb-impl-2.2.3-1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jcip-annotations-1.0-1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jersey-core-1.19.jar;D:\hadoop-3.2.1\share\h
adoop\common\lib\jersey-json-1.19.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jersey-server-1.19.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jersey-servlet-1.19.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jettis
on-1.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-http-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-io-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-security-9.3
.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-server-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-servlet-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jett
y-util-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-webapp-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jetty-xml-9.3.24.v20180605.jar;D:\hadoop-3.2.1\share\hadoop\common\li
b\jsch-0.1.54.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\json-smart-2.3.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jsp-api-2.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jsr305-3.0.0.jar;D:\hadoop-3.2.1\share
\hadoop\common\lib\jsr311-api-1.1.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\jul-to-slf4j-1.7.25.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-admin-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-c
lient-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-common-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-core-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-crypto-1.0.1.jar;D:\hadoop-3
.2.1\share\hadoop\common\lib\kerb-identity-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-server-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\common\lib\kerb-simplekdc-1.0.1.jar;D:\hadoop-3.2.1\share\hadoop\co
```

2. Starting dfs and yarn

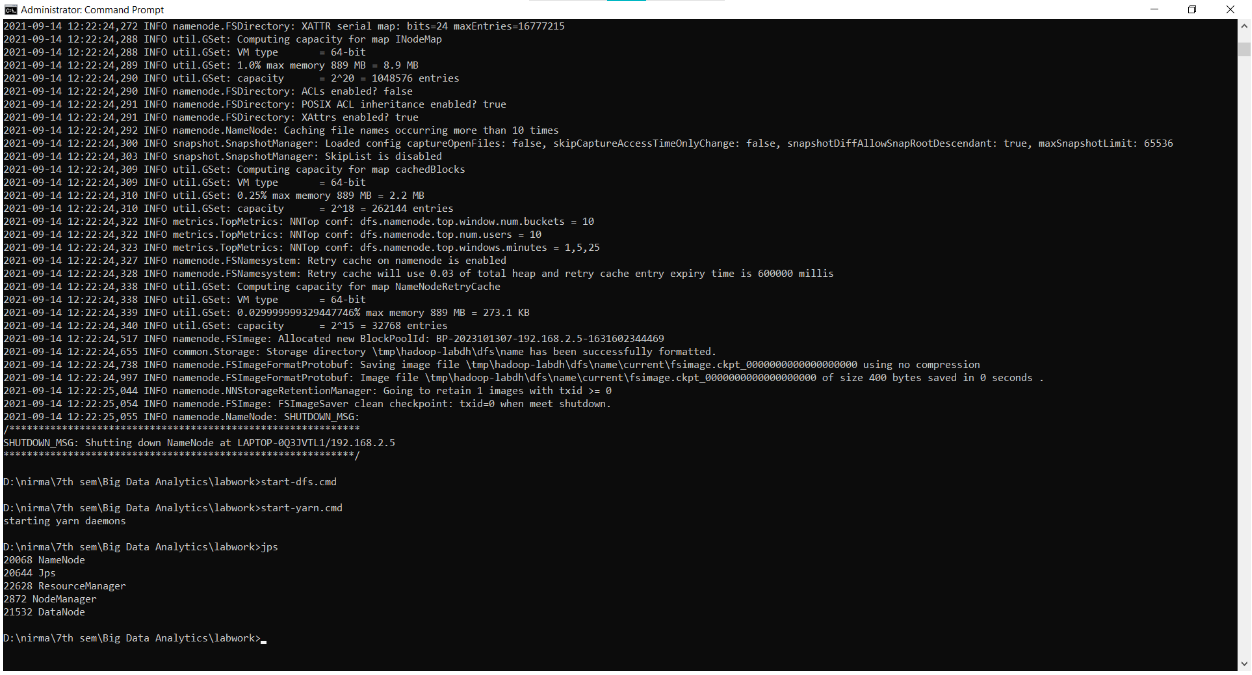

Administrator: Command Prompt

```
2021-09-14 12:22:24,272 INFO namenode.FSDirectory: XATTR serial map: bits=24 maxEntries=16777215
2021-09-14 12:22:24,288 INFO util.GSet: Computing capacity for map INodeMap
2021-09-14 12:22:24,288 INFO util.GSet: VM type       = 64-bit
2021-09-14 12:22:24,289 INFO util.GSet: 1.0% max memory 889 MB = 8.9 MB
2021-09-14 12:22:24,290 INFO util.GSet: capacity      = 2^20 = 1048576 entries
2021-09-14 12:22:24,290 INFO namenode.FSDirectory: ACLs enabled? false
2021-09-14 12:22:24,291 INFO namenode.FSDirectory: POSIX ACL inheritance enabled? true
2021-09-14 12:22:24,291 INFO namenode.FSDirectory: XAttrs enabled? true
2021-09-14 12:22:24,292 INFO namenode.NameNode: Caching file names occurring more than 10 times
2021-09-14 12:22:24,300 INFO snapshot.SnapshotManager: Loaded config captureOpenFiles: false, skipCaptureAccessTimeOnlyChange: false, snapshotDiffAllowSnapRootDescendant: true, maxSnapshotLimit: 65536
2021-09-14 12:22:24,303 INFO snapshot.SnapshotManager: SkipList is disabled
2021-09-14 12:22:24,309 INFO util.GSet: Computing capacity for map cachedBlocks
2021-09-14 12:22:24,309 INFO util.GSet: VM type       = 64-bit
2021-09-14 12:22:24,310 INFO util.GSet: 0.25% max memory 889 MB = 2.2 MB
2021-09-14 12:22:24,310 INFO util.GSet: capacity      = 2^18 = 262144 entries
2021-09-14 12:22:24,322 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.window.num.buckets = 10
2021-09-14 12:22:24,322 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.num.users = 10
2021-09-14 12:22:24,323 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.windows.minutes = 1,5,25
2021-09-14 12:22:24,327 INFO namenode.FSNamesystem: Retry cache on namenode is enabled
2021-09-14 12:22:24,328 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 600000 millis
2021-09-14 12:22:24,338 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2021-09-14 12:22:24,338 INFO util.GSet: VM type       = 64-bit
2021-09-14 12:22:24,339 INFO util.GSet: 0.029999999329447746% max memory 889 MB = 273.1 KB
2021-09-14 12:22:24,340 INFO util.GSet: capacity      = 2^15 = 32768 entries
2021-09-14 12:22:24,517 INFO namenode.FSImage: Allocated new BlockPoolId: BP-2023101307-192.168.2.5-1631602344469
2021-09-14 12:22:24,655 INFO common.Storage: Storage directory \tmp\hadoop-labdh\dfs\name has been successfully formatted.
2021-09-14 12:22:24,738 INFO namenode.FSImageFormatProtobuf: Saving image file \tmp\hadoop-labdh\dfs\name\current\fsimage.ckpt_0000000000000000000 using no compression
2021-09-14 12:22:24,997 INFO namenode.FSImageFormatProtobuf: Image file \tmp\hadoop-labdh\dfs\name\current\fsimage.ckpt_0000000000000000000 of size 400 bytes saved in 0 seconds .
2021-09-14 12:22:25,044 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2021-09-14 12:22:25,054 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2021-09-14 12:22:25,055 INFO namenode.NameNode: SHUTDOWN_MSG:
/************************************************************
SHUTDOWN_MSG: Shutting down NameNode at LAPTOP-0Q3JVTL1/192.168.2.5
************************************************************/

D:\nirma\7th sem\Big Data Analytics\labwork>start-dfs.cmd

D:\nirma\7th sem\Big Data Analytics\labwork>start-yarn.cmd
starting yarn daemons

D:\nirma\7th sem\Big Data Analytics\labwork>jps
20068 NameNode
20644 Jps
22628 ResourceManager
2872 NodeManager
21532 DataNode

D:\nirma\7th sem\Big Data Analytics\labwork>_
```

3. Checking the clusters and working on localhost

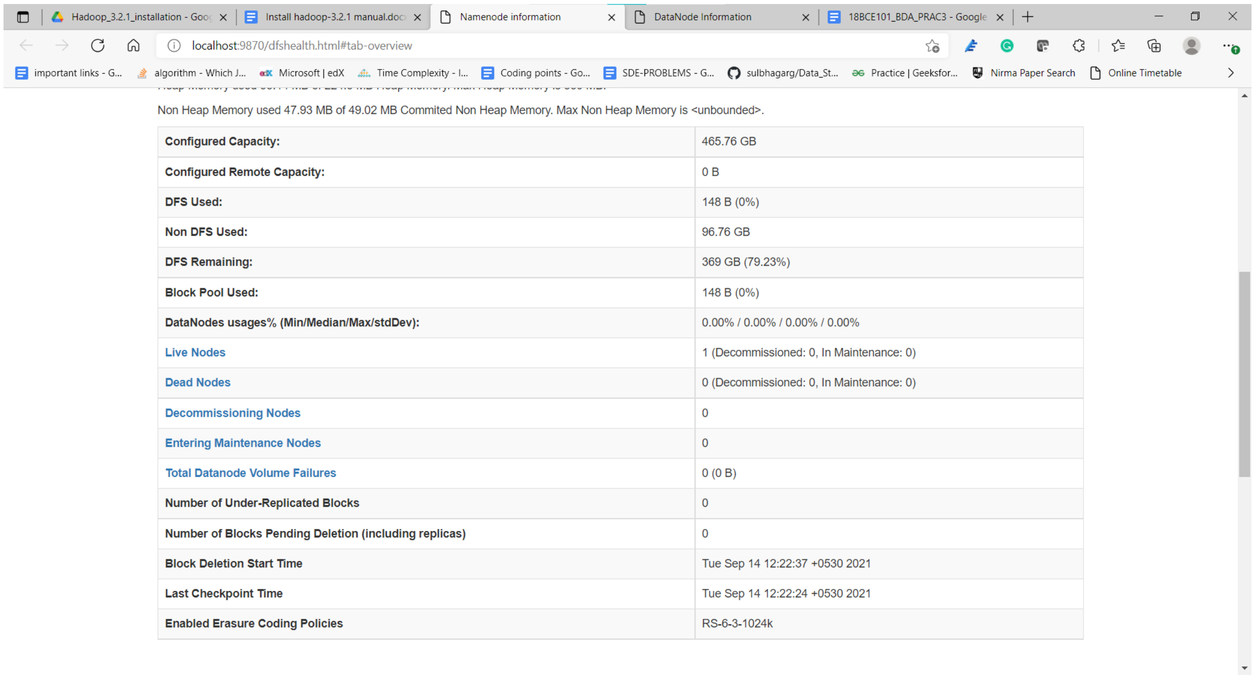

Non Heap Memory used 47.93 MB of 49.02 MB Commited Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	465.76 GB
Configured Remote Capacity:	0 B
DFS Used:	148 B (0%)
Non DFS Used:	96.76 GB
DFS Remaining:	369 GB (79.23%)
Block Pool Used:	148 B (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	1 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion (including replicas)	0
Block Deletion Start Time	Tue Sep 14 12:22:37 +0530 2021
Last Checkpoint Time	Tue Sep 14 12:22:24 +0530 2021
Enabled Erasure Coding Policies	RS-6-3-1024k

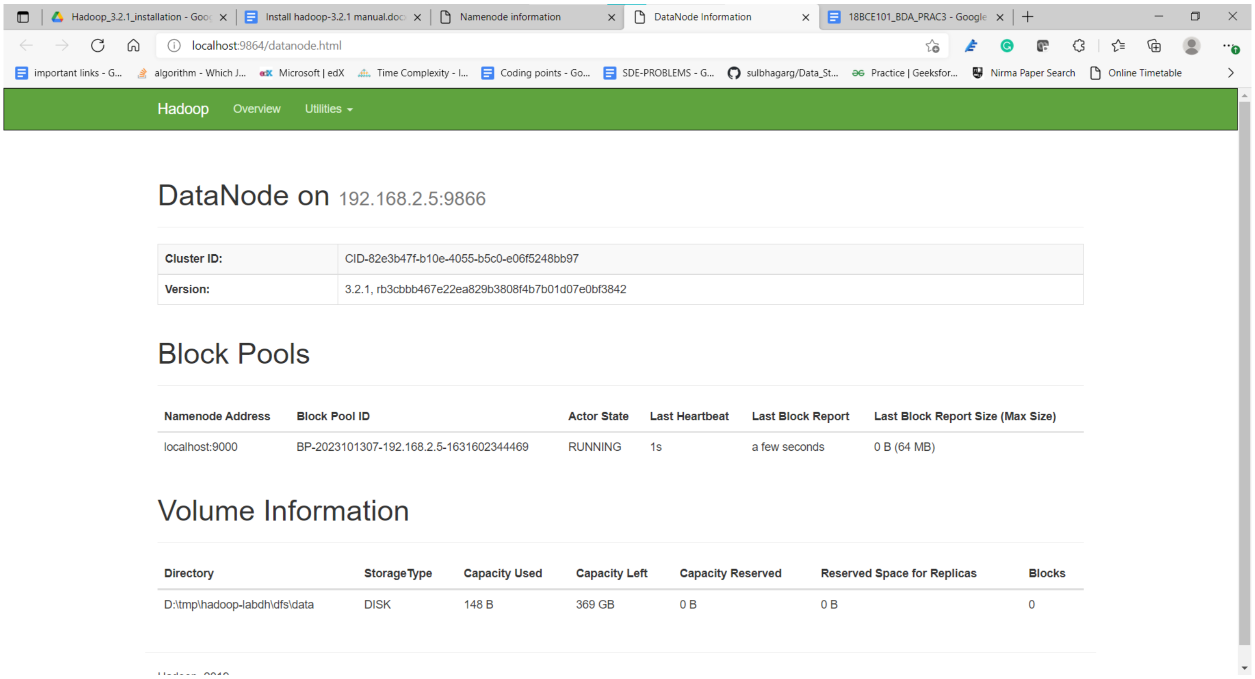

DataNode on 192.168.2.5:9866

Cluster ID:	CID-82e3b47f-b10e-4055-b5c0-e06f5248bb97
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842

Block Pools

Namenode Address	Block Pool ID	Actor State	Last Heartbeat	Last Block Report	Last Block Report Size (Max Size)
localhost:9000	BP-2023101307-192.168.2.5-1631602344469	RUNNING	1s	a few seconds	0 B (64 MB)

Volume Information

Directory	StorageType	Capacity Used	Capacity Left	Capacity Reserved	Reserved Space for Replicas	Blocks
D:\tmp\hadoop-labdh\dfs\data	DISK	148 B	369 GB	0 B	0 B	0

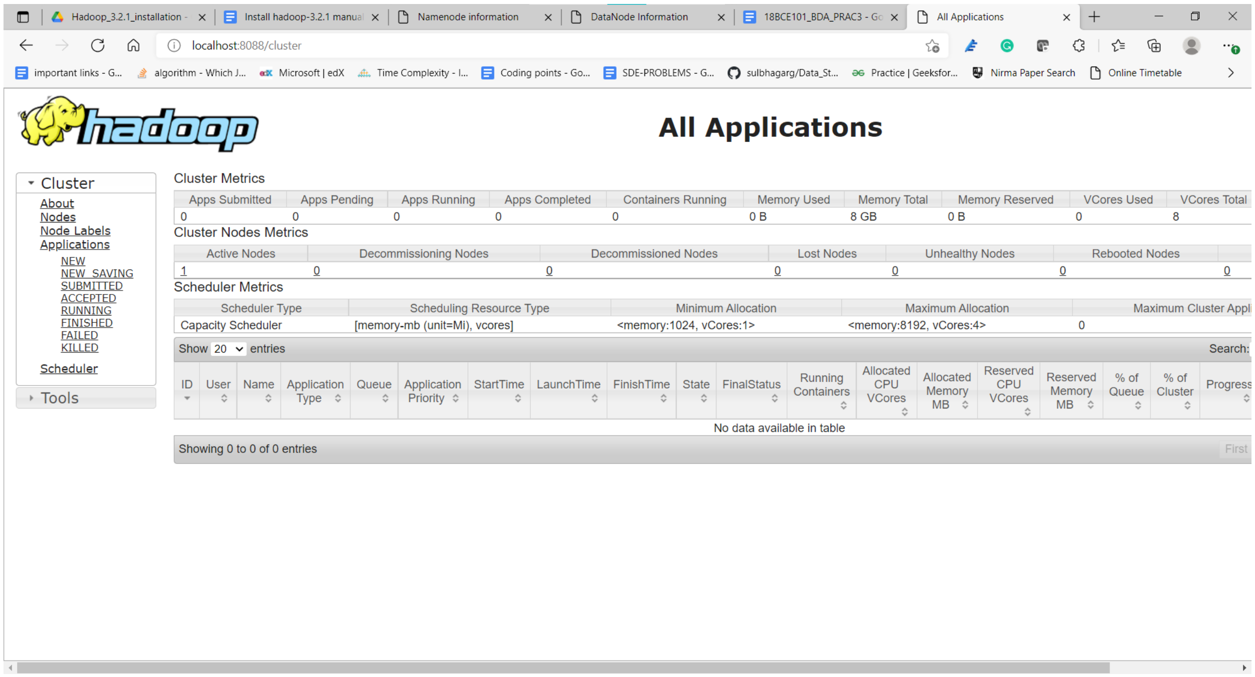

All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total
0	0	0	0	0	0 B	8 GB	0 B	0	8

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	
1	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Appl
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries   Search:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCores	Allocated Memory MB	Reserved CPU VCores	Reserved Memory MB	% of Queue	% of Cluster	Progress
No data available in table																		

Showing 0 to 0 of 0 entries   First

4. Mkdir command

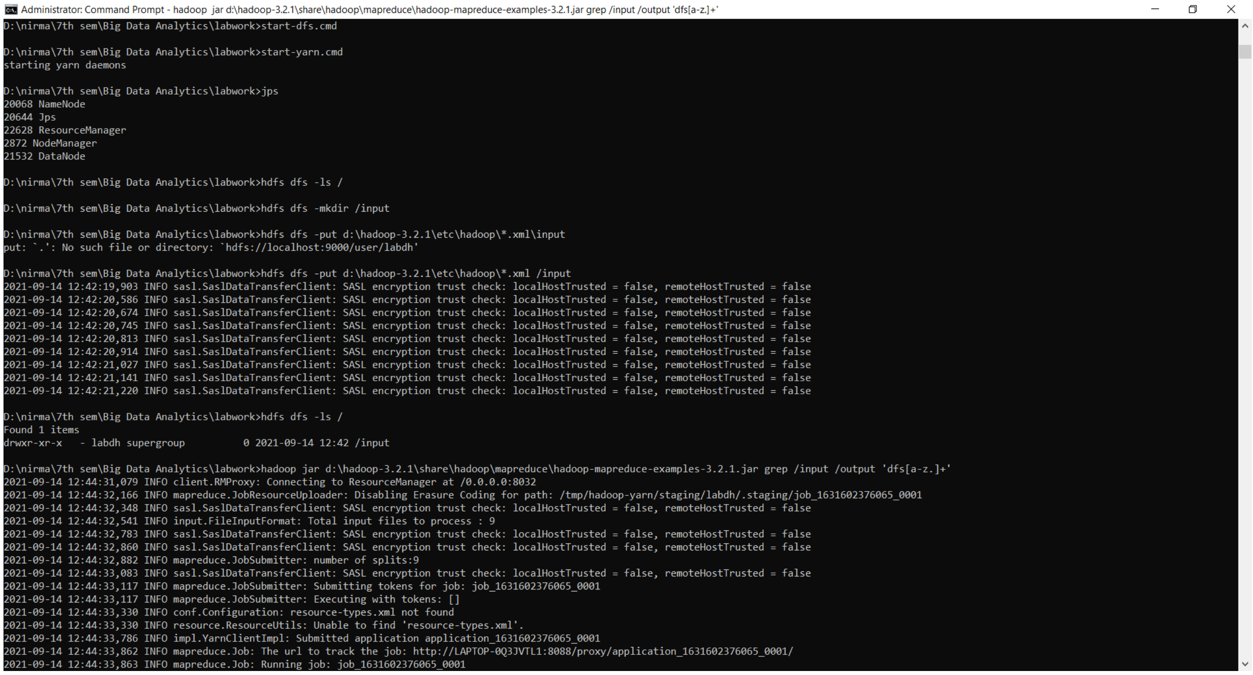

Administrator: Command Prompt - hadoop jar d:\hadoop-3.2.1\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.1.jar grep /input /output 'dfs[a-z.]+'

```
D:\nirma\7th sem\Big Data Analytics\labwork>start-dfs.cmd

D:\nirma\7th sem\Big Data Analytics\labwork>start-yarn.cmd
starting yarn daemons

D:\nirma\7th sem\Big Data Analytics\labwork>jps
20068 NameNode
20644 Jps
22628 ResourceManager
2872 NodeManager
21532 DataNode

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -ls /

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -mkdir /input

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -put d:\hadoop-3.2.1\etc\hadoop\*.xml\input
put: `.': No such file or directory: `hdfs://localhost:9000/user/labdh'

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -put d:\hadoop-3.2.1\etc\hadoop\*.xml /input
2021-09-14 12:42:19,903 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:20,586 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:20,674 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:20,745 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:20,813 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:20,914 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:21,027 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:21,141 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:42:21,220 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -ls /
Found 1 items
drwxr-xr-x   - labdh supergroup          0 2021-09-14 12:42 /input

D:\nirma\7th sem\Big Data Analytics\labwork>hadoop jar d:\hadoop-3.2.1\share\hadoop\mapreduce\hadoop-mapreduce-examples-3.2.1.jar grep /input /output 'dfs[a-z.]+'
2021-09-14 12:44:31,079 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
2021-09-14 12:44:32,166 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/labdh/.staging/job_1631602376065_0001
2021-09-14 12:44:32,348 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:44:32,541 INFO input.FileInputFormat: Total input files to process : 9
2021-09-14 12:44:32,783 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:44:32,860 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:44:32,882 INFO mapreduce.JobSubmitter: number of splits:9
2021-09-14 12:44:33,083 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2021-09-14 12:44:33,117 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1631602376065_0001
2021-09-14 12:44:33,117 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-09-14 12:44:33,330 INFO conf.Configuration: resource-types.xml not found
2021-09-14 12:44:33,330 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2021-09-14 12:44:33,786 INFO impl.YarnClientImpl: Submitted application application_1631602376065_0001
2021-09-14 12:44:33,862 INFO mapreduce.Job: The url to track the job: http://LAPTOP-0Q3JVTL1:8088/proxy/application_1631602376065_0001/
2021-09-14 12:44:33,863 INFO mapreduce.Job: Running job: job_1631602376065_0001
```

5. Grep command

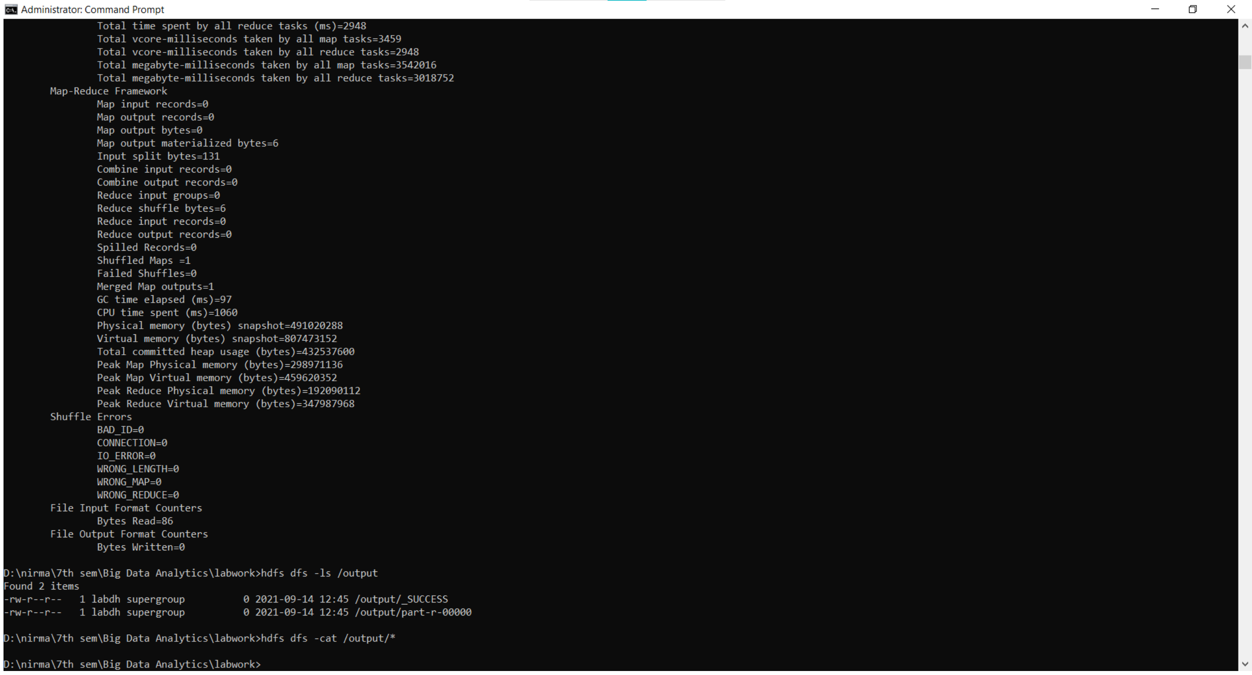

Administrator: Command Prompt

```
                Total time spent by all reduce tasks (ms)=2948
                Total vcore-milliseconds taken by all map tasks=3459
                Total vcore-milliseconds taken by all reduce tasks=2948
                Total megabyte-milliseconds taken by all map tasks=3542016
                Total megabyte-milliseconds taken by all reduce tasks=3018752
        Map-Reduce Framework
                Map input records=0
                Map output records=0
                Map output bytes=0
                Map output materialized bytes=6
                Input split bytes=131
                Combine input records=0
                Combine output records=0
                Reduce input groups=0
                Reduce shuffle bytes=6
                Reduce input records=0
                Reduce output records=0
                Spilled Records=0
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=97
                CPU time spent (ms)=1060
                Physical memory (bytes) snapshot=491020288
                Virtual memory (bytes) snapshot=807473152
                Total committed heap usage (bytes)=432537600
                Peak Map Physical memory (bytes)=298971136
                Peak Map Virtual memory (bytes)=459620352
                Peak Reduce Physical memory (bytes)=192090112
                Peak Reduce Virtual memory (bytes)=347987968
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=86
        File Output Format Counters
                Bytes Written=0

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -ls /output
Found 2 items
-rw-r--r--   1 labdh supergroup          0 2021-09-14 12:45 /output/_SUCCESS
-rw-r--r--   1 labdh supergroup          0 2021-09-14 12:45 /output/part-r-00000

D:\nirma\7th sem\Big Data Analytics\labwork>hdfs dfs -cat /output/*

D:\nirma\7th sem\Big Data Analytics\labwork>
```

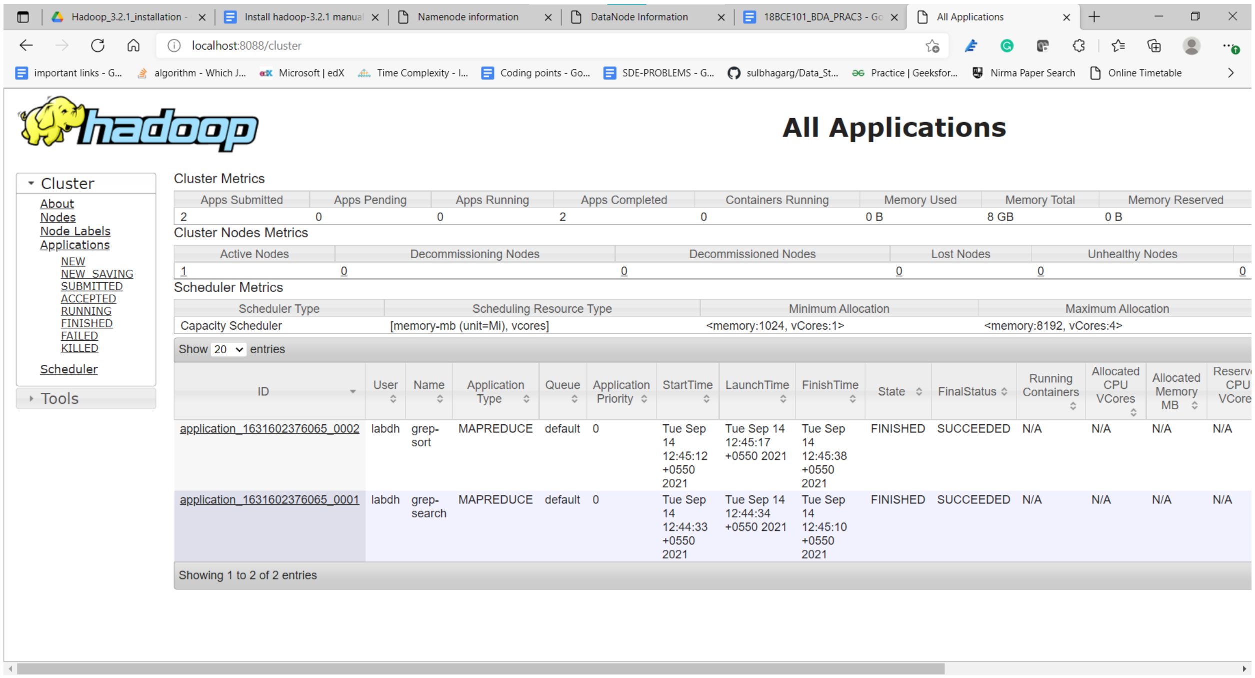

All Applications

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved
2	0	0	2	0	0 B	8 GB	0 B

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	
1	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>

Show 20 entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCores	Allocated Memory MB	Reserve CPU VCore
application_1631602376065_0002	labdh	grep-sort	MAPREDUCE	default	0	Tue Sep 14 12:45:12 +0550 2021	Tue Sep 14 12:45:17 +0550 2021	Tue Sep 14 12:45:38 +0550 2021	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A
application_1631602376065_0001	labdh	grep-search	MAPREDUCE	default	0	Tue Sep 14 12:44:33 +0550 2021	Tue Sep 14 12:44:34 +0550 2021	Tue Sep 14 12:45:10 +0550 2021	FINISHED	SUCCEEDED	N/A	N/A	N/A	N/A

Showing 1 to 2 of 2 entries